

Introduction to Machine Learning

Pak Alexander Alexandrovich

Candidate of tech. sciences

a.pak@kbtu.kz

“I keep saying that the sexy job in the next 10 years will be statisticians”

Hal Varian, Chief Economist Google

Lecture 0: Introduction

Course objective

- To learn the history of pattern recognition and machine learning algorithms.
- To learn and be able to choose a model to solve practical issues.
- To learn the techniques of preprocessing and postprocessing of data.
- To learn the design of numerical experiments, namely verification, validation and output analysis.
- To apply knowledge and skills learned on at least one large collaborative modeling project.

Assessment

Type of activity	Finalscores
Exercises, assignments	28
Quiz	8
SIS	8
Midterm	16
Final Exam + Project	40
Total	100%

History of Artificial Intelligence



- 1936-1956 The birth of theory
- 1956-1976 Golden Age
- 1969-1980 Crisis of ANN

The birth of AI



On computable numbers, with an app to
the Entscheidungsproblem (1936)
Computing machinery and intelligence
(1950)

Finite state machine or A-machine of A.Turing



The head is always over a particular square of the tape;



Finite state machine or A-machine of A.Turing


The behaviour of the computer at any moment is determined by the symbols which **he** is observing and **his** 'state of mind' at that moment,

Alan Turing

Example of calculation with the help of Turing Machine

23 <u>56</u> 8	23 <u>56</u> 138	23 <u>56</u> 138 115 1288
---	---	--

Turing Test



The Death of the Father of AI sci

Turing was prosecuted in 1952 for homosexual acts, when by the Labouchere Amendment, "gross indecency" was criminal in the UK. He accepted chemical castration treatment, with DES, as an alternative to prison. Turing died in 1954, 16 days before his 42nd birthday, from cyanide poisoning. An inquest determined his death as suicide, but it has been noted that the known evidence is also consistent with accidental poisoning.

A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY* (1943)




Warren
McCulloch



Walter Pitts

Formal Neuron of McCulloch and Pitts



Evolutional calculus



Nils Aall Barricelli
Genetic Algo (1954)

McCarthy et al. Proposal for the project

Мы предлагаем двухмесячный исследовательский семинар в составе десяти человек для исследования искусственного интеллекта в течение лета 1956 года в Дортмундском колледже Гановера, Нью-Хэмпшир. Отправной точкой исследования является убеждение в том, что все аспекты обучения, и других проявлений интеллекта, могут быть настолько точно описаны, что машина может запрограммирована на их выполнение. Будет сделана попытка выяснить, как машины могут использовать язык, делать абстракции, решать различные виды задач, которые пока решает лишь человек, и самообучаться. Мы полагаем, что возможно существенное продвижение в этом вопросе, если тщательно отобранная группа ученых будет совместно работать над ним в течение лета.

Golden Age 1956-1976



John Alan Robinson

Prolog language (1972)



Alan Colmerauer

Philippe Roussel



Lotfi Zadeh

Fuzzy Logics (1965)

Вес для O

4

X	O	X
X	O	
O		

2

4

Вес для X

1

2

2

2

2

1

X	O	X
X	O	X
O	O	

X	O	X
X	O	X
O		O

X	O	X
X	O	O
O	X	

X	O	X
X	O	
O	X	O

X	O	X
X	O	O
O		X

X	O	X
X	O	
O	O	X

O выигрывает

O выигрывает

Вес для O

2

2

2

2

X	O	X
X	O	X
O		O

X	O	X
X	O	O
O	X	

X	O	X
X	O	
O	X	O

X	O	X
X	O	O
O		X

Ничья

Ничья

Ничья

Ничья

Natural Language Processing



Joseph Weizenbaum

ELIZA — A Computer Program for the Study of Natural Language Communication
between Man and Machine (1966)



Neural Nets



Frank Rosenblatt

Principles of neurodynamics: Perceptrons and the theory of brain mechanisms (1962)

Evolutional calculus




John Henry Holland

Minsky Papert





Back Prop (1974)



1986

Big Data

- Widespread use of personal computers and wireless communication leads to “big data”
- We are both producers and consumers of data
- Data is not random, it has structure, e.g., customer behavior
- We need “big theory” to extract that structure from data for
 - (a) Understanding the process
 - (b) Making predictions for the future

Why “Learn” ?

26

- Machine learning is programming computers to optimize a performance criterion using example data or past experience.
- There is no need to “learn” to calculate payroll
- Learning is used when:
 - Human expertise does not exist (navigating on Mars),
 - Humans are unable to explain their expertise (speech recognition)
 - Solution changes in time (routing on a computer network)
 - Solution needs to be adapted to particular cases (user biometrics)

What We Talk About When We Talk About “Learning”

- Learning general models from a data of particular examples
- Data is cheap and abundant (data warehouses, data marts); knowledge is expensive and scarce.
- Example in retail: Customer transactions to consumer behavior:
People who bought “Blink” also bought “Outliers”
(www.amazon.com)
- Build a model that is *a good and useful approximation* to the data.

Data Mining

28

- **Retail:** Market basket analysis, Customer relationship management (CRM)
- **Finance:** Credit scoring, fraud detection
- **Manufacturing:** Control, robotics, troubleshooting
- **Medicine:** Medical diagnosis
- **Telecommunications:** Spam filters, intrusion detection
- **Bioinformatics:** Motifs, alignment
- **Web mining:** Search engines
- ...

What is Machine Learning?

29

- Optimize a performance criterion using example data or past experience.
- Role of Statistics: Inference from a sample
- Role of Computer science: Efficient algorithms to
 - Solve the optimization problem
 - Representing and evaluating the model for inference

Applications

30

- Association
- Supervised Learning
 - Classification
 - Regression
- Unsupervised Learning
- Reinforcement Learning

Learning Associations

31

- Basket analysis:


$P(Y | X)$ probability that somebody who buys X also buys Y where X and Y are products/services.

Example: $P(\text{chips} | \text{beer}) = 0.7$

Classification

32

- Example: Credit scoring
- Differentiating between **low-risk** and **high-risk** customers from their *income* and *savings*



Discriminant: IF $income > \theta_1$ AND $savings > \theta_2$
THEN **low-risk** ELSE **high-risk**

Classification: Applications

33

- Aka Pattern recognition
- **Face recognition:** Pose, lighting, occlusion (glasses, beard), make-up, hair style
- **Character recognition:** Different handwriting styles.
- **Speech recognition:** Temporal dependency.
- **Medical diagnosis:** From symptoms to illnesses
- **Biometrics:** Recognition/authentication using physical and/or behavioral characteristics: Face, iris, signature, etc
- **Outlier/novelty detection:**

Face Recognition

34

Training examples of a person



Test images



ORL dataset,
AT&T Laboratories, Cambridge UK

Regression

- Example: Price of a used car


• x : car attributes

y : price

$$y = g(x \mid \theta)$$

$g(\cdot)$ model,


θ parameters




Regression Applications

36

- Navigating a car: Angle of the steering
- Kinematics of a robot arm



■ Response surface design



Supervised Learning: Uses

37

- **Prediction of future cases:** Use the rule to predict the output for future inputs
- **Knowledge extraction:** The rule is easy to understand
- **Compression:** The rule is simpler than the data it explains
- **Outlier detection:** Exceptions that are not covered by the rule, e.g., fraud

Unsupervised Learning

38

- Learning “what normally happens”
- No output
- Clustering: Grouping similar instances
- Example applications
 - Customer segmentation in CRM
 - Image compression: Color quantization
 - Bioinformatics: Learning motifs

Reinforcement Learning

39

- Learning a policy: A **sequence** of outputs
- No supervised output but delayed reward
- Credit assignment problem
- Game playing
- Robot in a maze
- Multiple agents, partial observability, ...

Resources: Datasets

40

- UCI Repository:
<http://www.ics.uci.edu/~mlearn/MLRepository.html>
- Statlib: <http://lib.stat.cmu.edu/>
- Kaggle: <http://www.Kaggle.com/>

Datasets

country	incomeperperson	alcoholconsumption	armedforcesrate	breastcancerper100th
Afghanistan		0.03	0.5696534	26.8
Albania	1914.996551	7.29	1.0247361	57.4
Algeria	2231.993335	0.69	2.306817	23.5
Andorra	21943.3399	10.17		
Angola	1381.004268	5.57	1.4613288	23.1

Codebook

Variable name: The name or number assigned to each variable in the data collection. Some researchers prefer to use mnemonic abbreviations (e.g., EMPLOY1), while others use alphanumeric patterns (e.g., VAR001). For survey data, try to name variables after the question numbers - e.g., Q1, Q2b, etc. [In above example, H40-SF12-2]

Variable label: A brief description to identify the variable for the user. Where possible, use the exact question or research wording. ["SF12 - ASSESSMENT OF R'S GENERAL HEALTH"]

Question text: Where applicable, the exact wording from survey questions. ["In general, would you say your health is . . ."]

Values: The actual coded values in the data for this variable. [1, 2, 3, 4, 5]

Value labels: The textual descriptions of the codes. [Excellent, Very Good, Good, Fair, Poor]

Codebook

Summary statistics: Where appropriate and depending on the type of variable, provide unweighted summary statistics for quick reference. For categorical variables, for instance, frequency counts showing the number of times a value occurs and the percentage of cases that value represents for the variable are appropriate. For continuous variables, minimum, maximum, and median values are relevant.

Missing data: Where applicable, the values and labels of missing data. Missing data can bias an analysis and is important to convey in study documentation. Remember to describe all missing codes, including "system missing" and blank. [e.g., Refusal (-1)]

Universe skip patterns: Where applicable, information about the population to which the variable refers, as well as the preceding and following variables. [e.g., Default Next Question: H00035.00]

Notes: Additional notes, remarks, or comments that contextualize the information conveyed in the variable or relay special instructions. For measures or questions from copyrighted instruments, the notes field is the appropriate location to cite the source.

Codebook

Variable Name	Description of Indicator	Main Source
incomeperperson	2010 Gross Domestic Product per capita in constant 2000 US\$. The inflation but not the differences in the cost of living between countries has been taken into account.	World Bank Work Development Indicators
alconsumption	2008 alcohol consumption per adult (age 15+), litres Recorded and estimated average alcohol consumption, adult (15+) per capita consumption in litres pure alcohol	WHO
armedforcesrate	Armed forces personnel (% of total labor force)	Work Development Indicators
breastcancerper100TH	2002 breast cancer new cases per 100,000 female Number of new cases of breast cancer in 100,000 female residents during the certain year.	ARC (International Agency for Research on Cancer)
co2emissions	2006 cumulative CO2 emission (metric tons), Total amount of CO2 emission in metric tons since 1751.	CDIAC (Carbon Dioxide Information Analysis Center)

Variable: Categorical

- The blood type of a person: A, B, AB or O.
- The state that a person lives in.
- The political party that a voter in a European country might vote for: Christian Democrat, Social Democrat, Green Party, etc.
- The type of a rock: igneous, sedimentary or metamorphic.
- The identity of a particular word (e.g., in a language model): One of V possible choices, for a vocabulary of size V .

Variable: Quantitative

Variable Measurement

- Height Inches, feet, centimeters,
- Temperature Celsius, Fahrenheit, Kelvin, Réaumur...
- Age Years, months, decades, minutes
- Weight Pounds, tons, ounces, grams
- Area Acres, square miles, square feet
- Speed Miles per hour, light years, feet per second

Example with NESARC

3649-3649	ALCABDEPP12DX	ALCOHOL ABUSE/DEPENDENCE PRIOR TO THE LAST 12 MONTHS
	31677	0. No alcohol diagnosis
	6994	1. Alcohol abuse only
	563	2. Alcohol dependence only
	3859	3. Alcohol abuse and dependence
3650-3650	TAB12MDX	NICOTINE DEPENDENCE IN THE LAST 12 MONTHS
	38131	0. No nicotine dependence
	4962	1. Nicotine dependence
3652-3652	TABLIFEDX	NICOTINE DEPENDENCE - LIFETIME
	36156	0. No nicotine dependence
	6937	1. Nicotine dependence
3653-3653	STIM12ABDEP	AMPHETAMINE ABUSE/DEPENDENCE IN LAST 12 MONTHS
	43032	0. No amphetamine diagnosis
	34	1. Amphetamine abuse only
	3	2. Amphetamine dependence only
	24	3. Amphetamine abuse and dependence

Resources: Journals

48

- Journal of Machine Learning Research www.jmlr.org
- Machine Learning
- Neural Computation
- Neural Networks
- IEEE Trans on Neural Networks and Learning Systems
- IEEE Trans on Pattern Analysis and Machine Intelligence
- Journals on Statistics/Data Mining/Signal Processing/Natural Language Processing/Bioinformatics/...

Resources: Conferences

49

- International Conference on Machine Learning (ICML)
- European Conference on Machine Learning (ECML)
- Neural Information Processing Systems (NIPS)
- Uncertainty in Artificial Intelligence (UAI)
- Computational Learning Theory (COLT)
- International Conference on Artificial Neural Networks (ICANN)
- International Conference on AI & Statistics (AISTATS)
- International Conference on Pattern Recognition (ICPR)
- ...